

# Attentive Graph Neural Networks

Prof. O-Joun Lee



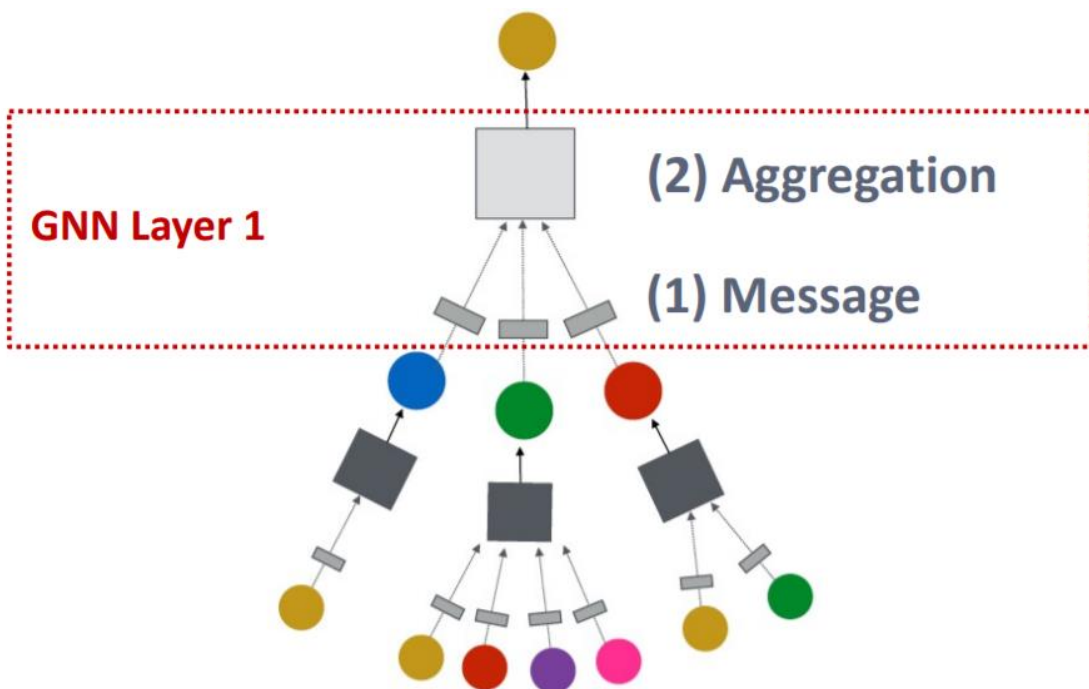
Dept. of Artificial Intelligence,  
The Catholic University of Korea  
*ojlee@catholic.ac.kr*

# Contents



- Graph neural network issues.
- Attention in Graph neural networks.
- Attention in Heterogeneous graphs.
- GAT sample code.


- GNN Layer = **Message** + **Aggregation**
  - Message COMPUTATION
    - how to make each neighborhood node as embedding?
  - Message AGGERGATION
    - how to combine those embeddings?



Update rule: 
$$\mathbf{h}_i^{(l+1)} = \sigma \left( \mathbf{h}_i^{(l)} \mathbf{W}_0^{(l)} + \sum_{j \in \mathcal{N}_i} \frac{1}{c_{ij}} \mathbf{h}_j^{(l)} \mathbf{W}_1^{(l)} \right)$$

- **Intuition:** Each node will create a message, which will be sent to other nodes later
- **Example:** A Linear layer  $\mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}$ 
  - Multiply node features with weight matrix  $\mathbf{W}^{(l)}$

Message function:  $\mathbf{m}_u^{(l)} = \text{MSG}^{(l)} \left( \mathbf{h}_u^{(l-1)} \right)$







- **Intuition:** Each node will aggregate the messages from node  $v$ 's neighbors

$$\mathbf{h}_v^{(l)} = \text{AGG}^{(l)} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right)$$

- **Example:** Sum( $\cdot$ ), Mean( $\cdot$ ) or Max( $\cdot$ ) aggregator

$$\mathbf{h}_v^{(l)} = \text{Sum}(\{\mathbf{m}_u^{(l)}, u \in N(v)\})$$



➤ **Issue:** Information from node  $v$  itself could get lost

➤ Computation of  $\mathbf{h}_v^{(l)}$  does not directly depend on  $\mathbf{h}_v^{(l-1)}$

➤ **Solution:** Include  $\mathbf{h}_v^{(l-1)}$  when computing  $\mathbf{h}_v^{(l)}$

➤ (1) **Message:** compute message from node  $v$  itself

$$\text{●} \text{●} \text{●} \quad \mathbf{m}_u^{(l)} = \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \qquad \text{●} \quad \mathbf{m}_v^{(l)} = \mathbf{B}^{(l)} \mathbf{h}_v^{(l-1)}$$

➤ (2) **Aggregation:** After aggregating from neighbors, we can aggregate the message from node  $v$  itself

➤ Via concatenation or summation

$$\mathbf{h}_v^{(l)} = \text{CONCAT} \left( \underbrace{\text{AGG} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right)}_{\text{First aggregate from neighbors}}, \underbrace{\mathbf{m}_v^{(l)}}_{\text{Then aggregate from node itself}} \right)$$


## ➤ Pure Graph Convolutional Networks (GCN)

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{w}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} \right)$$

## ➤ **Message:** Each neighbour $u$ :

$$\mathbf{m}_u^{(l)} = \frac{1}{|N(v)|} \mathbf{w}^{(l)} \mathbf{h}_u^{(l-1)}$$

equally important to  $v$




## ➤ **Aggregation:** Sum over messages from neighbors, then apply activation

$$\mathbf{h}_v^{(l)} = \sigma \left( \text{Sum} \left( \left\{ \mathbf{m}_u^{(l)}, u \in N(v) \right\} \right) \right)$$

→ All neighbors  $u \in N(v)$  are equally important to node  $v$


## Not all node's neighbors are equally important

- Attention is a mechanism that allows a network to focus on certain parts of the input when processing it
- The attention focuses on the important parts of the input data and fades out the rest.
  - **Idea:** the neural network should devote more computing power on that small but important part of the data.
  - Which part of the data is more important depends on the context and is learned through training.



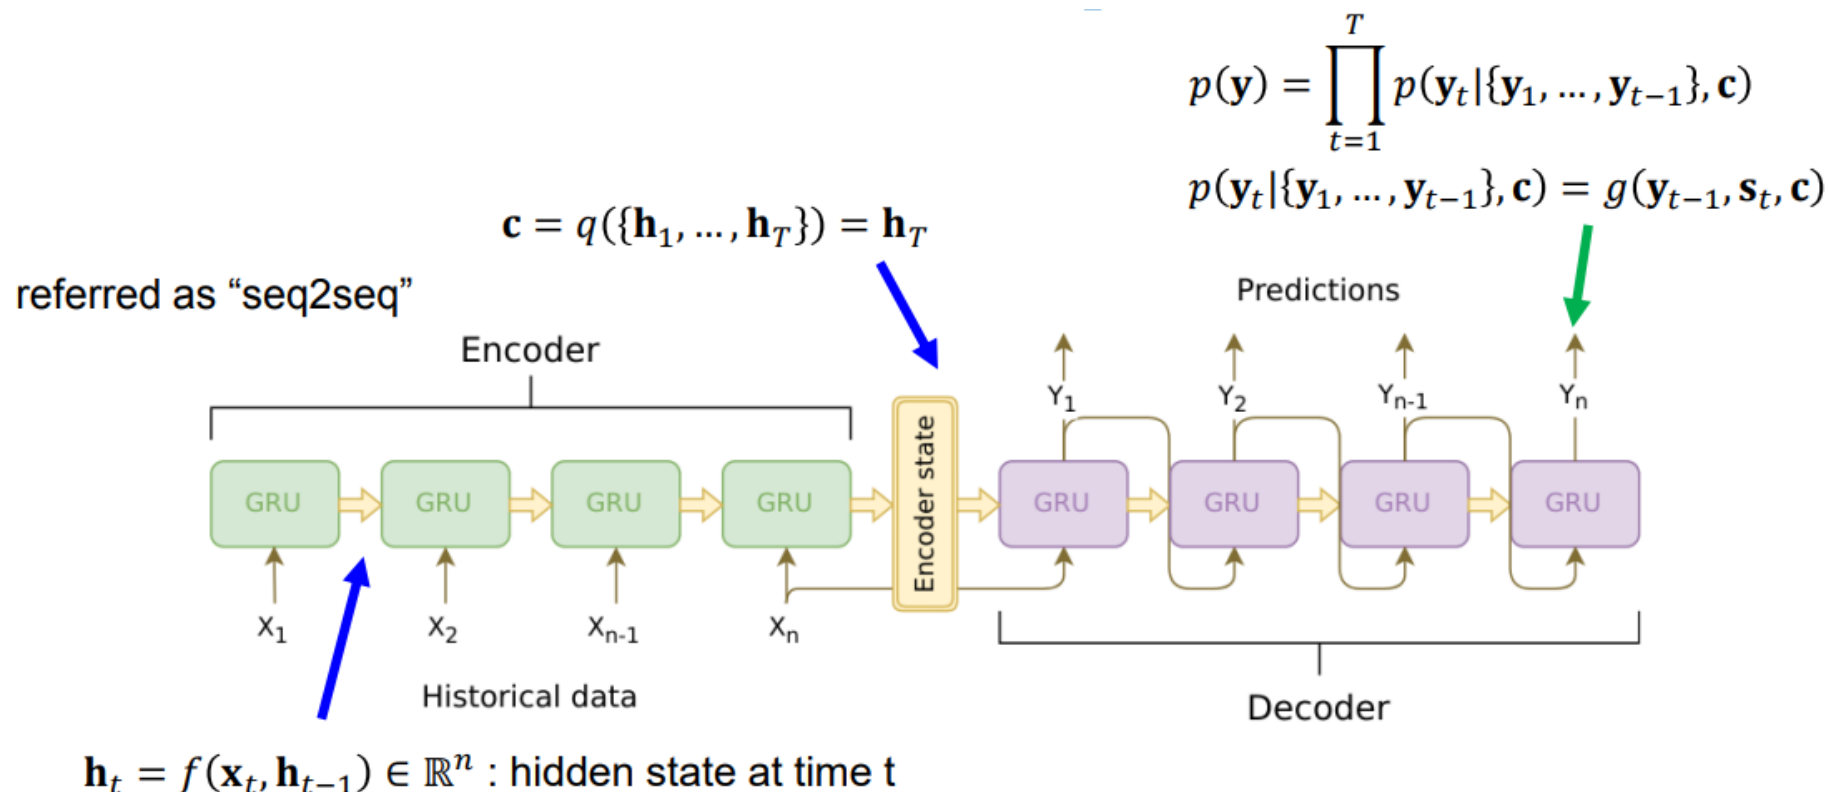


- We deduce something by paying attention to something that is relatively more important.




## ➤ RNN encoder-decoder for neural machine translation:


- In capability of remembering long sentences : Often it has forgotten the first part once it completes processing the whole input. The attention mechanism was born to resolve this problem.



- GNN compute node representations from representations of neighbours.
- Nodes can have largely different neighbourhood sizes.
- Not all neighbours have relevant information for a certain node.
- Attention mechanism allow to adaptively weight the contribution of each neighbour when updating a node.




- **Attention means:** assign an attention coefficient to each neighbor, indicating the importance of that neighbor's features for the feature update of the node.




[Figure from Veličković et al. (ICLR 2018)]

- The key difference between GAT and GCN is how the information from the one-hop neighborhood is aggregated.


Vanilla GCN updates information of neighbor nodes with same importance



Attention mechanism enables GCN to update nodes with different importance.



- In Graph Attention Networks (GATs), the concept of multiple attention heads is similar to the idea of multiple filters in Convolutional Neural Networks (CNNs).
- Each attention head can potentially learn to pay attention to different types of neighborhood information.






- Input node features: Each node in the graph has a feature vector.

$$\mathbf{h} = \{\vec{h}_1, \vec{h}_2, \dots, \vec{h}_N\}, \vec{h}_i \in \mathbb{R}^F$$

- Calculate energy (co-efficient) between two nodes

$$e_{ij} = a(\mathbf{W}\vec{h}_i, \mathbf{W}\vec{h}_j)$$

a: attention function



- Attention score (over the neighbors): Normalize over all the neighbors

$$\alpha_{ij} = \frac{\exp \left( \text{LeakyReLU} \left( \vec{\mathbf{a}}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_j] \right) \right)}{\sum_{k \in \mathcal{N}_i} \exp \left( \text{LeakyReLU} \left( \vec{\mathbf{a}}^T [\mathbf{W} \vec{h}_i \| \mathbf{W} \vec{h}_k] \right) \right)}$$


- Multi-head attention

- Feature concatenation

$$\vec{h}'_i = \parallel_{k=1}^K \sigma \left( \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$

- Feature averaging (for the final layer)

$$\vec{h}'_i = \sigma \left( \frac{1}{K} \sum_{k=1}^K \sum_{j \in \mathcal{N}_i} \alpha_{ij}^k \mathbf{W}^k \vec{h}_j \right)$$




## ➤ Pros:

- No need to score intermediate edge-based activation vectors (when using dot product attention).
- Slower than GCNs but faster than GNNs with edge embeddings.

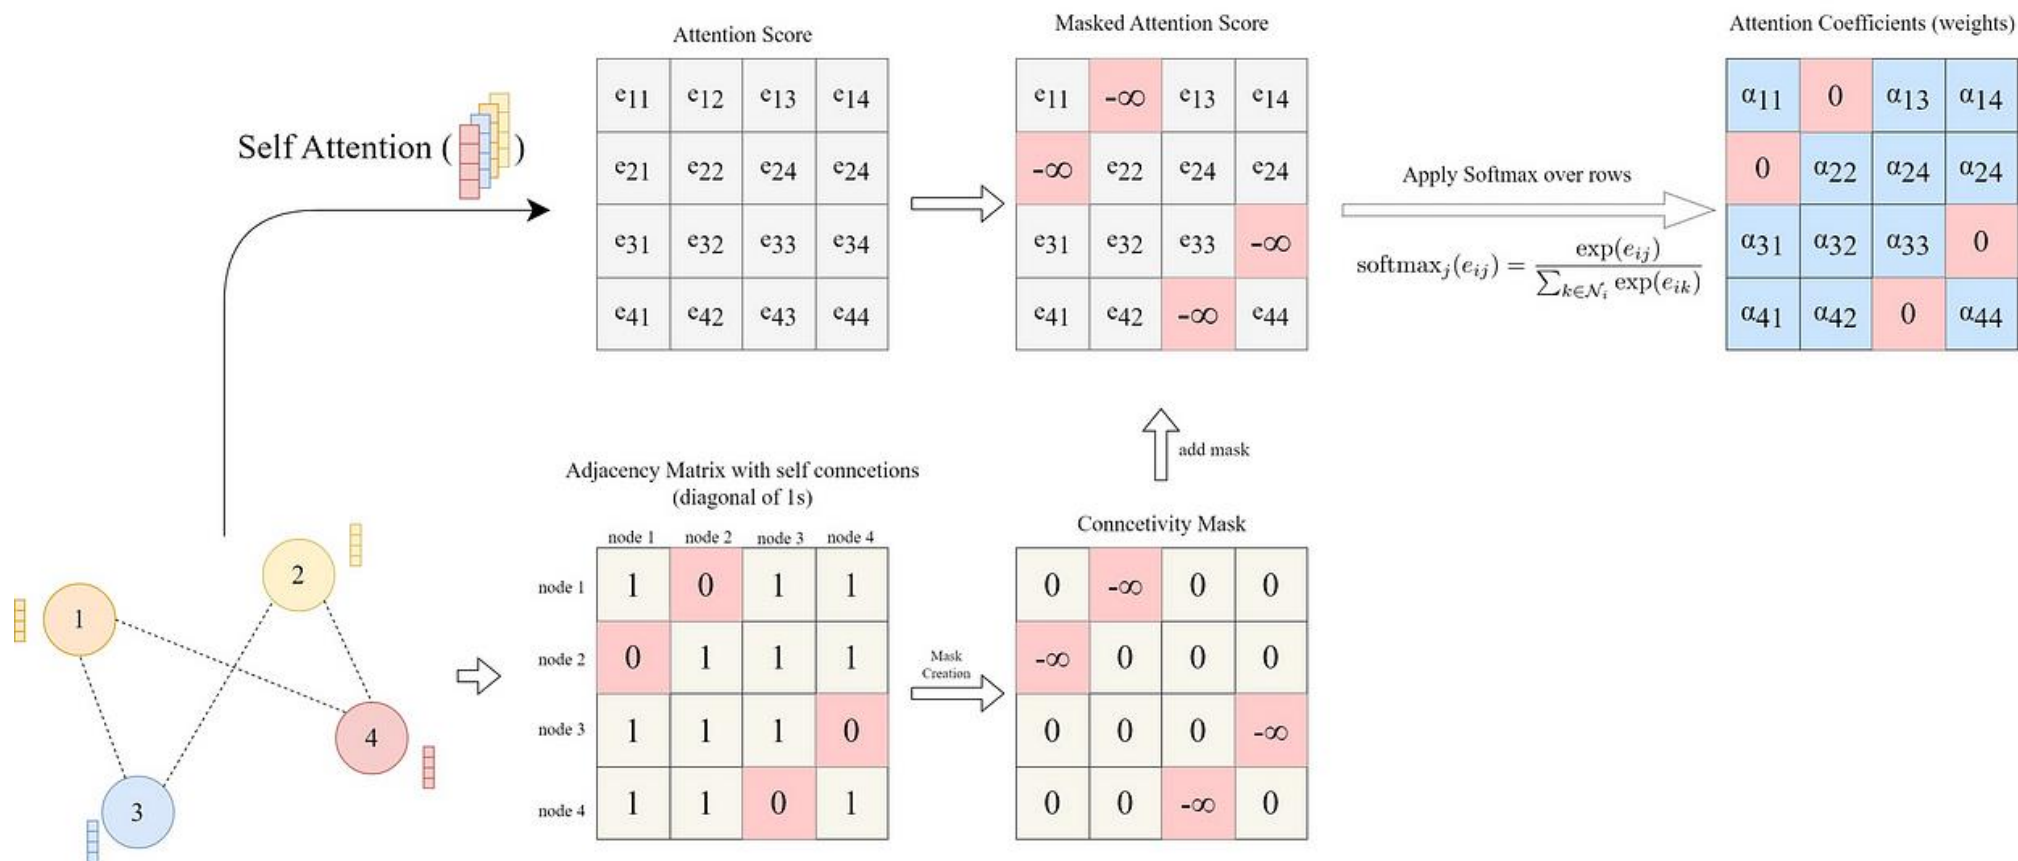
## ➤ Cons:

- Can be more difficult to optimize.


- The whole operation is illustrated below:



- Applying masking mechanism to the masked attention score, then apply Softmax function:



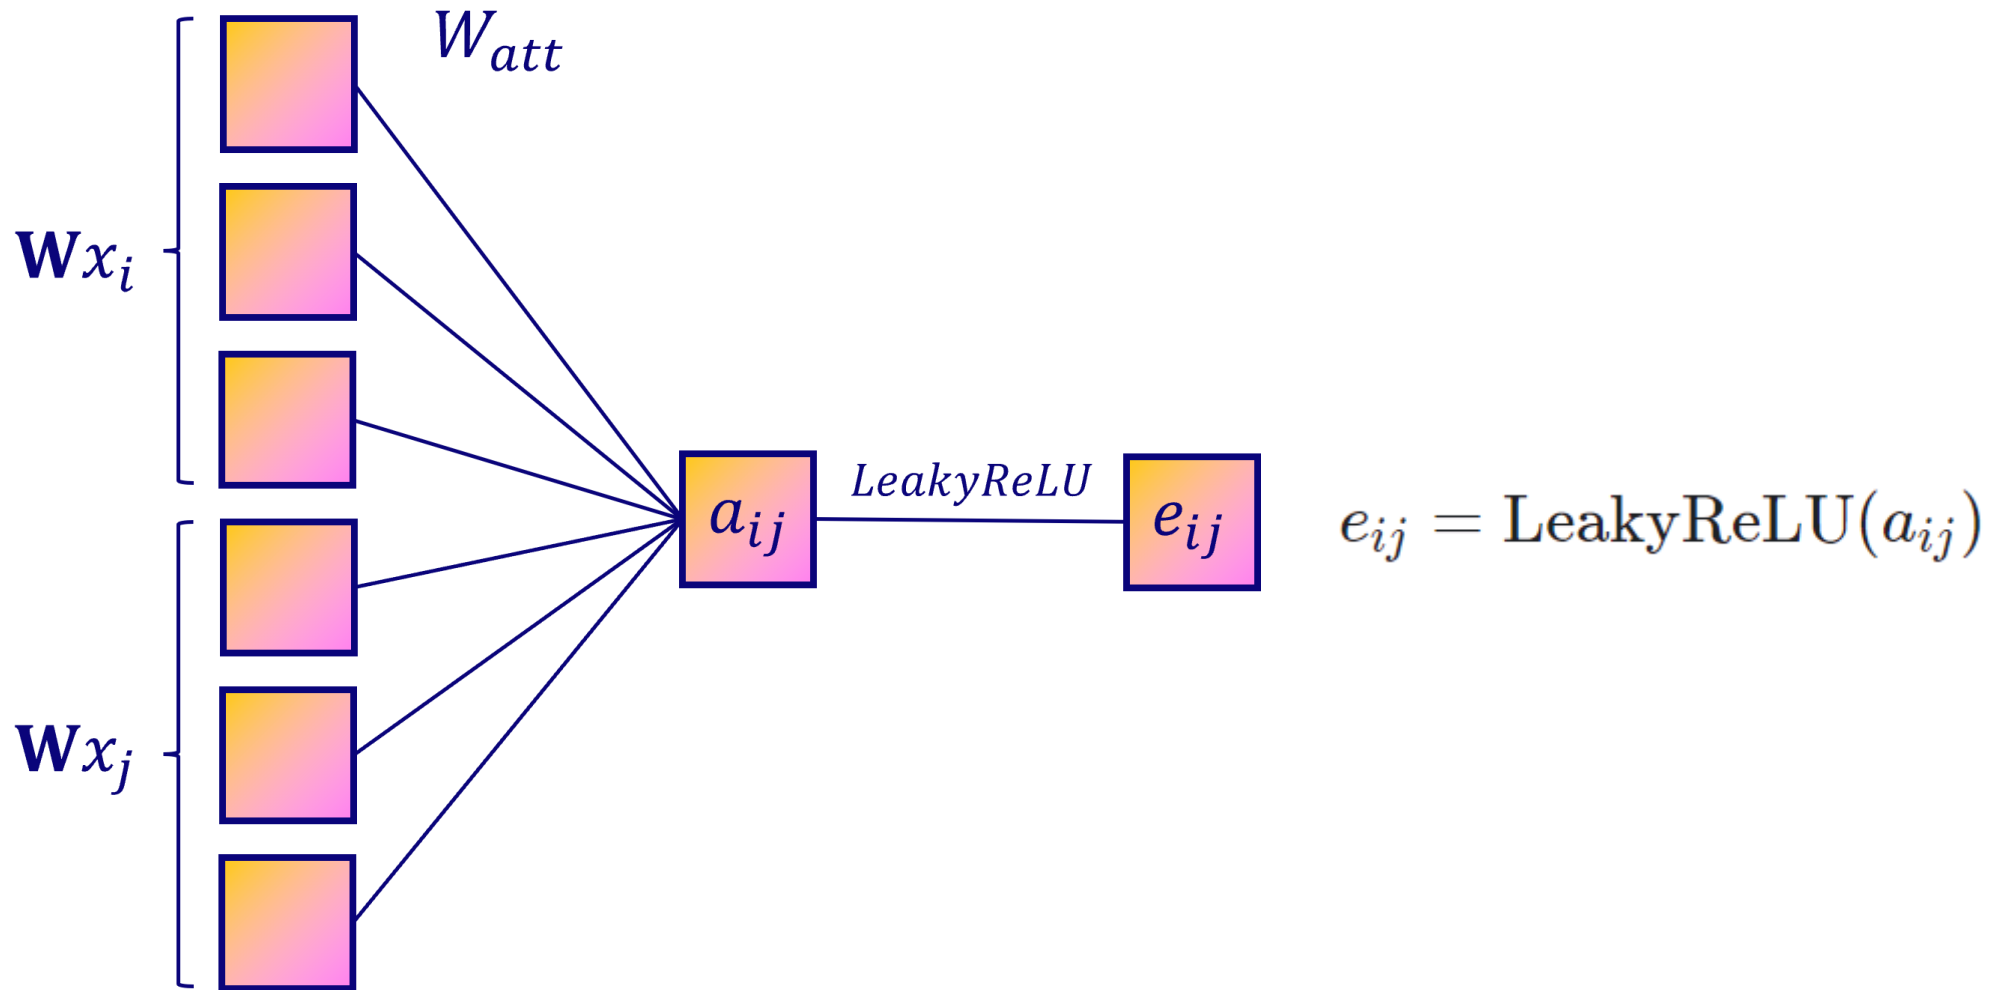
- **Linear transformation:** To calculate the attention coefficient, we need to consider pairs of nodes. An easy way to create these pairs is to concatenate attribute vectors from both nodes.



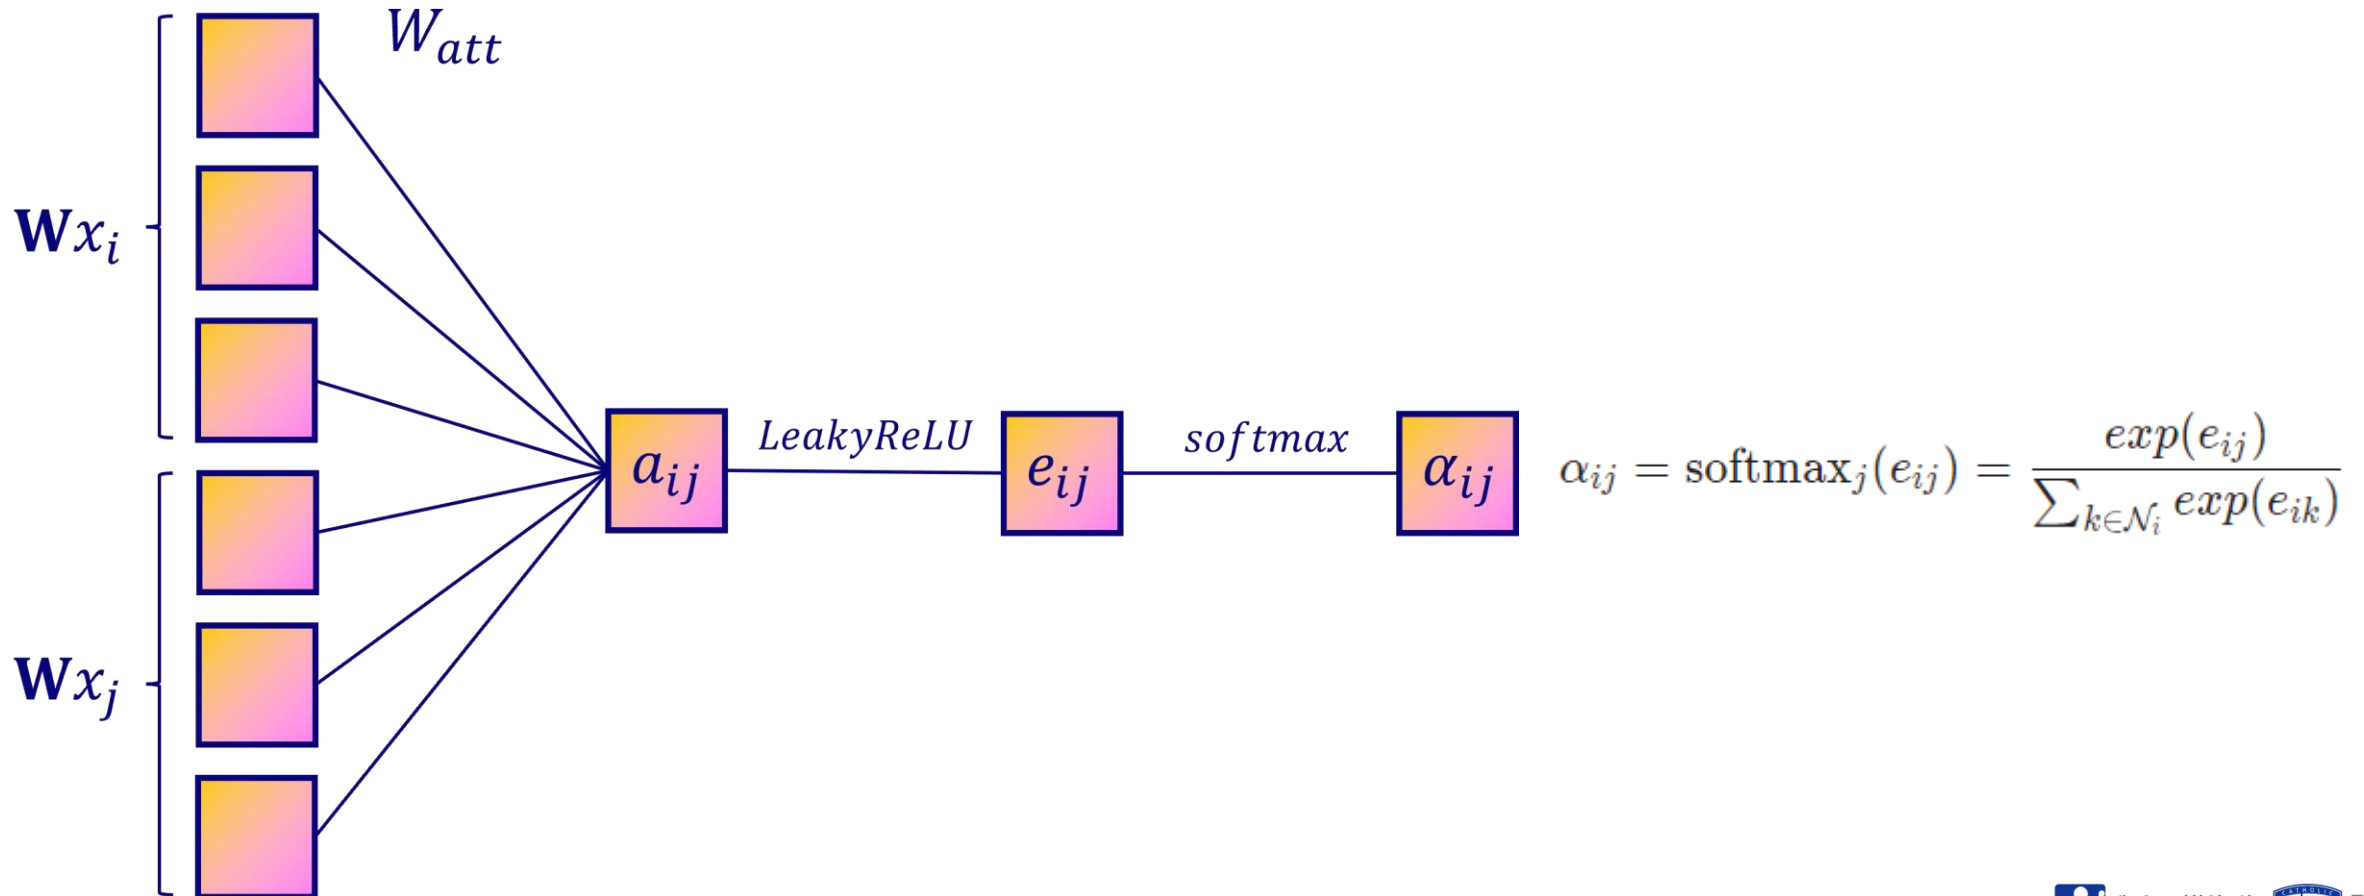
$$a_{ij} = W_{att}^t [\mathbf{W}x_i \parallel \mathbf{W}x_j]$$



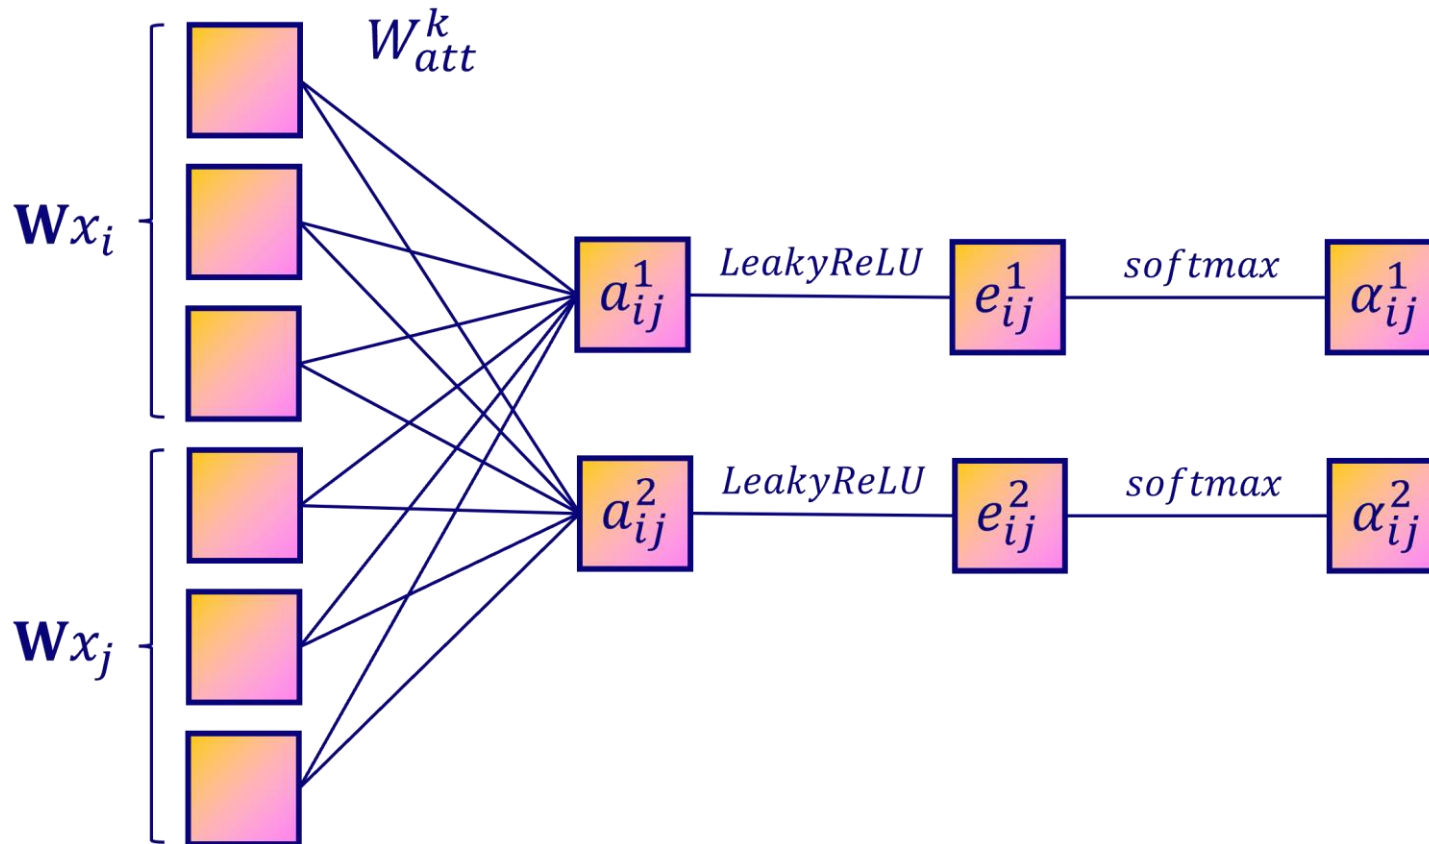
- **Activation function:** add nonlinearity with an activation function. In this case, the paper's authors chose the LeakyReLU function.



- Softmax normalization: The output of our neural network is not normalized, which is a problem since we want to compare these coefficients.
- A common way to do it with neural networks is to use the softmax function.



- Multi-head attention: In GATs, multi-head attention consists of replicating the same three steps several times in order to average or concatenate the results.



**Average:**

$$h_i = \frac{1}{n} \sum_{k=1}^n h_i^k$$

**Concatenation:**

$$h_i = \parallel_{k=1}^n h_i^k$$

🏠 / torch\_geometric.nn / conv.GATConv

## conv.GATConv

```
class GATConv ( in_channels: Union[int, Tuple[int, int]], out_channels: int, heads: int = 1, concat: bool =
True, negative_slope: float = 0.2, dropout: float = 0.0, add_self_loops: bool = True, edge_dim:
Optional[int] = None, fill_value: Union[float, Tensor, str] = 'mean', bias: bool = True, **kwargs )
[source]
```

Bases: MessagePassing

The graph attentional operator from the "Graph Attention Networks" paper

$$\mathbf{x}'_i = \alpha_{i,i} \Theta \mathbf{x}_i + \sum_{j \in \mathcal{N}(i)} \alpha_{i,j} \Theta \mathbf{x}_j,$$

where the attention coefficients  $\alpha_{i,j}$  are computed as

$$\alpha_{i,j} = \frac{\exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_j]))}{\sum_{k \in \mathcal{N}(i) \cup \{i\}} \exp(\text{LeakyReLU}(\mathbf{a}^\top [\Theta \mathbf{x}_i \parallel \Theta \mathbf{x}_k]))}.$$

```
46  class GAT(torch.nn.Module):
47      def __init__(self, num_features, num_classes, dims, drop=0.0):
48          super(GAT, self).__init__()
49          heads = 8
50          self.conv1 = GATConv(num_features, dims, heads=heads, dropout=0.3, concat=False)
51          # On the Pubmed dataset, use heads=8 in conv2.
52          self.conv2 = GATConv(dims, num_classes, heads=heads, concat=False,
53                               dropout=0.3)
54          self.drop = torch.nn.Dropout(p=drop)
55      def forward(self, x, edge_index):
56          x = F.elu(self.conv1(x, edge_index))
57          x = self.drop(x)
58          x = self.conv2(x, edge_index)
59          return F.log_softmax(x, dim=1), x
```

➤ Let's try some simple GAT code in the sample code file

- GATv2s is similar to GAT.
- The GATv2 operator fixes the static attention problem of the standard GAT.
  - Static attention is when the attention to the key nodes has the same rank (order) for any query node.
  - GAT computes attention from query node  $i$  to key node  $j$ :

$$\begin{aligned} e_{ij} &= \text{LeakyReLU}\left(\mathbf{a}^\top \left[ \mathbf{W} \vec{h}_i \parallel \mathbf{W} \vec{h}_j \right]\right) \\ &= \text{LeakyReLU}\left(\mathbf{a}_1^\top \mathbf{W} \vec{h}_i + \mathbf{a}_2^\top \mathbf{W} \vec{h}_j\right) \end{aligned}$$


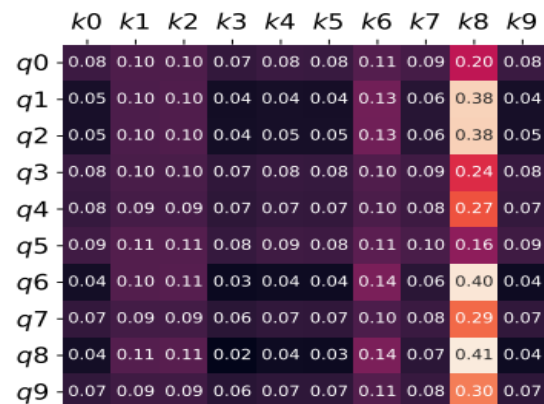
GAT

$$\begin{aligned} e_{ij} &= \mathbf{a}^\top \text{LeakyReLU}\left(\mathbf{W} \left[ \vec{h}_i \parallel \vec{h}_j \right]\right) \\ &= \mathbf{a}^\top \text{LeakyReLU}\left(\mathbf{W}_l \vec{h}_i + \mathbf{W}_r \vec{h}_j\right) \end{aligned}$$



GATv2



- The GATv2 model performs better than the first version GAT, because it uses a dynamic graph attention variant that has a universal approximator attention function, it is more expressive than the other model, based on a static attention



Attention in standard GAT



Attention in GATv2

- GATv2 is available as part of PyTorch Geometric library

```
from torch_geometric.nn import GATv2Conv
```

[🏠](#) / [torch\\_geometric.nn](#) / [conv.GATv2Conv](#)

## conv.GATv2Conv

```
class GATv2Conv ( in_channels: Union[int, Tuple[int, int]], out_channels: int, heads: int = 1, concat: bool = True, negative_slope: float = 0.2, dropout: float = 0.0, add_self_loops: bool = True, edge_dim: Optional[int] = None, fill_value: Union[float, Tensor, str] = 'mean', bias: bool = True, share_weights: bool = False, **kwargs ) \[source\]
```

Bases: `MessagePassing`

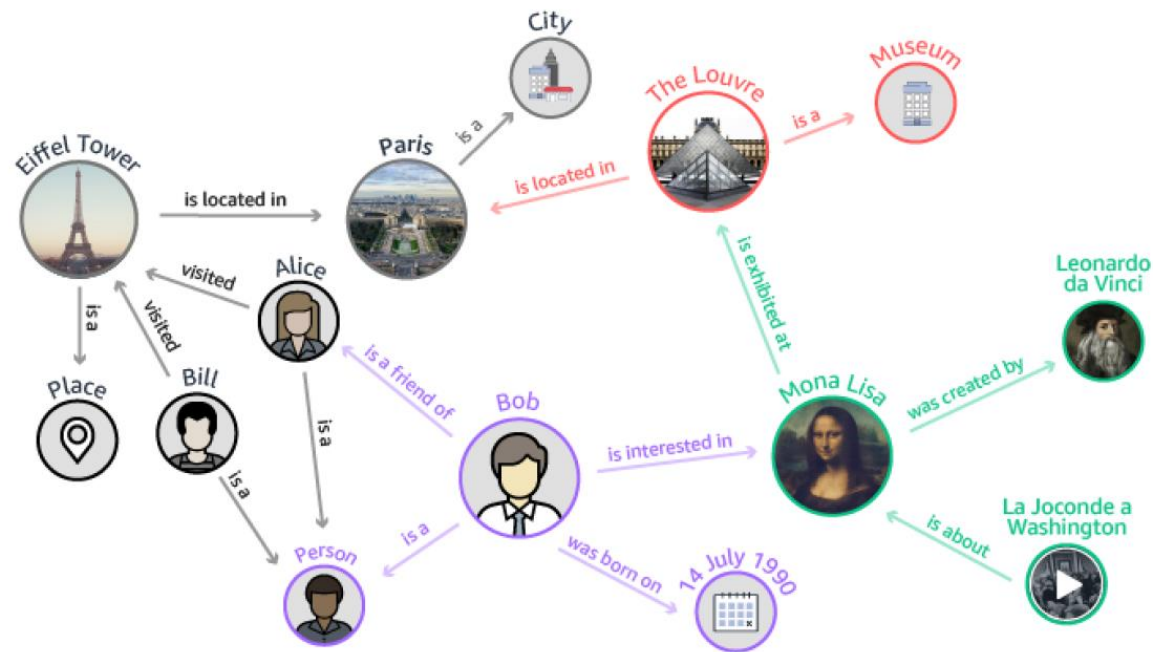
The GATv2 operator from the “How Attentive are Graph Attention Networks?” paper, which fixes the static attention problem of the standard `GATConv` layer. Since the linear layers in the standard GAT are applied right after each other, the ranking of attended nodes is unconditioned on the query node. In contrast, in `GATv2`, every node can attend to any other node.

- GATv2 is available as part of PyTorch Geometric library


```
from torch_geometric.nn import GATv2Conv
```

```
136 class GATv2(torch.nn.Module):
137     def __init__(self, num_features, num_classes, dims, drop=0.0):
138         super(GATv2, self).__init__()
139         heads = 8
140         self.h = None
141         self.conv1 = GATv2Conv(num_features, dims, heads=heads, dropout = 0.3, concat=False)
142         self.conv2 = GATv2Conv(dims, num_classes, heads=heads, concat=False, dropout=0.3)
143         self.drop = torch.nn.Dropout(p=drop)
144     def forward(self, x, edge_index, g, Kindices):
145         x = F.elu(self.conv1(x, edge_index))
146         x = self.drop(x)
147         x = self.conv2(x, edge_index)
148         self.h = x
149         return F.log_softmax(x, dim=1)
```


- Graph in real world:
  - Many node types, link types
  - Non-ordered
  - Complex connections



- Multiple types of nodes or links
- Rich semantic information
  - Meta-path: a relation sequence connecting objects  
(e.g., Movie-Actor-Movie).




Homogeneous Graph



(a) Node Type (b) Heterogeneous Graph

Heterogeneous Graph



Movie-Director-Movie

Two movies directed by the same director.




Movie-Actor-Movie

Two movies are starred by the same actor.

## ➤ DBLP Bibliographic network

- Node (type)
  - KDD (Venue)
  - Author 1
- Link (Type)
  - Write ( Author - Paper)
  - Publish ( Paper – Venue)

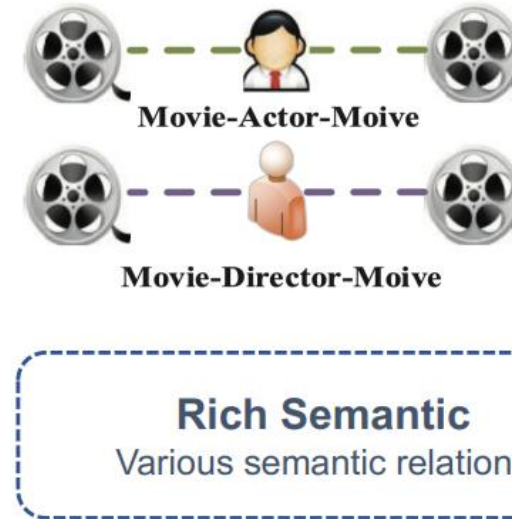



A. Examples of A-P-V-P-A meta-path on DBLP

B. Examples of common neighborhood objects between two authors in DBLP

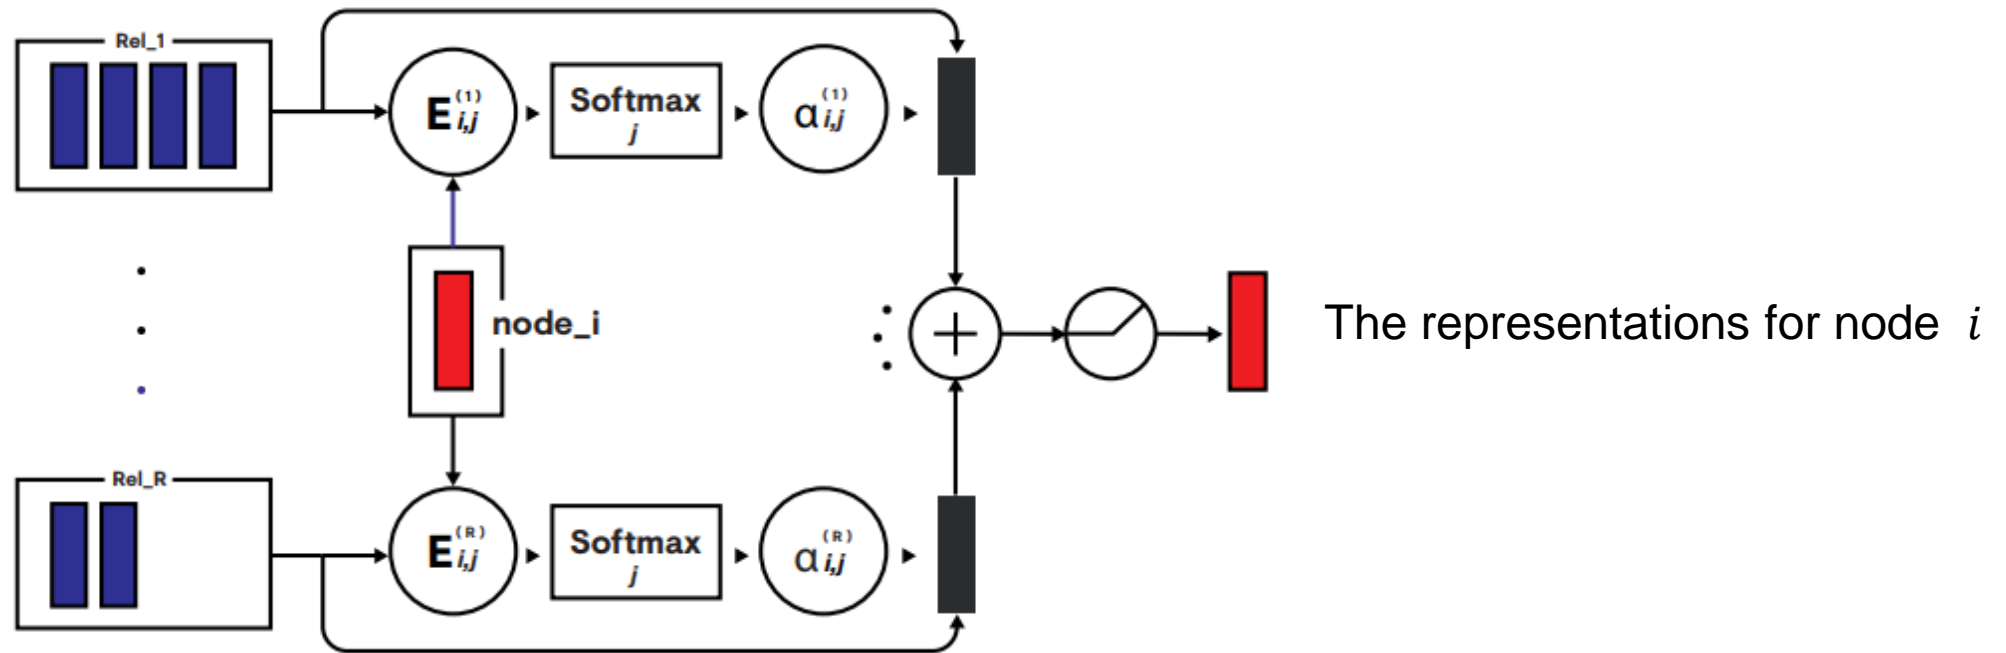


- Existing GNNs focus on homogeneous graphs
  - Cannot handle multiple types of nodes and edges.
  - Cannot capture rich semantic information.



**Challenge:** How to handle the heterogeneity of graph?

- The objective: Extending attention mechanisms to the relational graph domain



A target node  $i$  have different relations :  $Rel_1, Rel_2, \dots, Rel_R$

The logits  $E_{i,j}^{(r)}$  of each relation  $r$ :  $E_{i,j}^{(r)} = a \left( g_i^{(r)}, g_j^{(r)} \right)$ ,

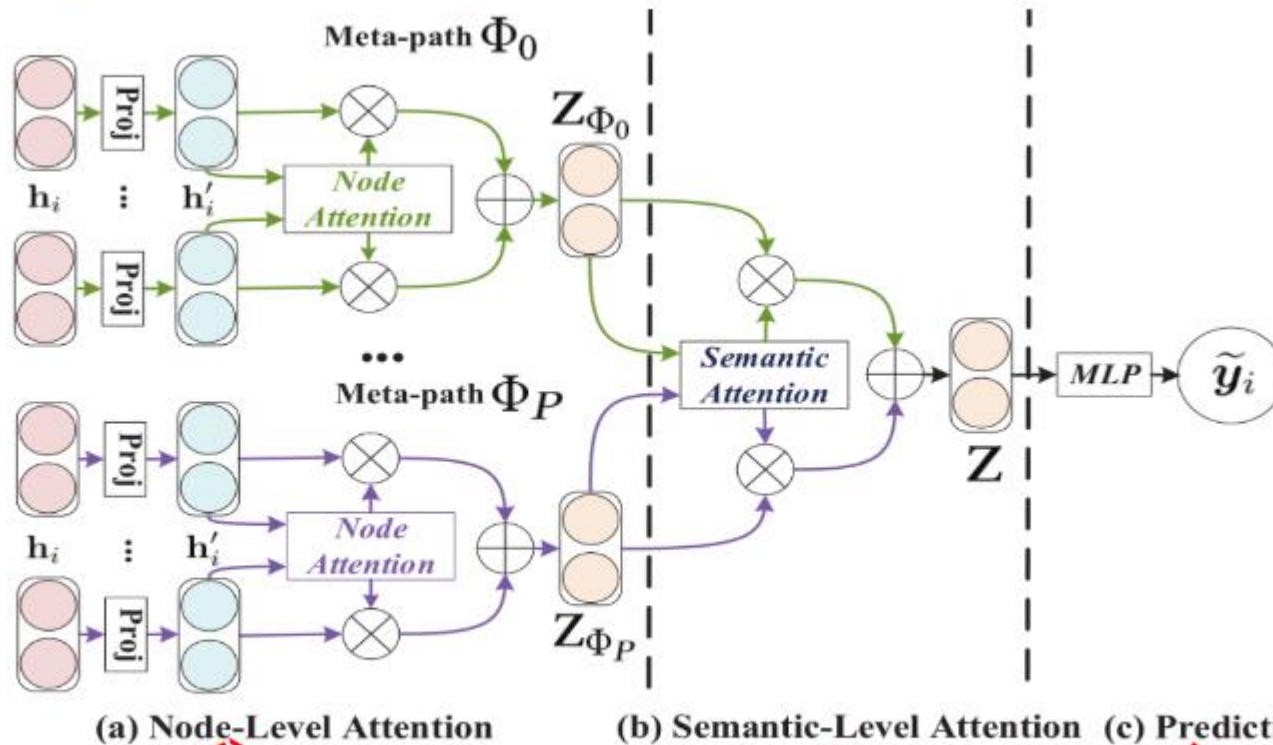
where:  $G^{(r)} = H W^{(r)} \in \mathbb{R}^{N \times F'}$ , the representation feature matrix under relation  $r$

- RGAT is available as part of PyTorch Geometric library

```
from torch_geometric.nn import RGATConv
```

```
16  ✓ class RGAT(torch.nn.Module):  
17  ✓     def __init__(self, in_channels, hidden_channels, out_channels,  
18             num_relations):  
19             super().__init__()  
20             self.conv1 = RGATConv(in_channels, hidden_channels, num_relations)  
21             self.conv2 = RGATConv(hidden_channels, hidden_channels, num_relations)  
22             self.lin = torch.nn.Linear(hidden_channels, out_channels)  
23  
24  ✓     def forward(self, x, edge_index, edge_type):  
25             x = self.conv1(x, edge_index, edge_type).relu()  
26             x = self.conv2(x, edge_index, edge_type).relu()  
27             x = self.lin(x)  
28             return F.log_softmax(x, dim=-1)
```

## Heterogeneous Graph Attention Network (HAN)


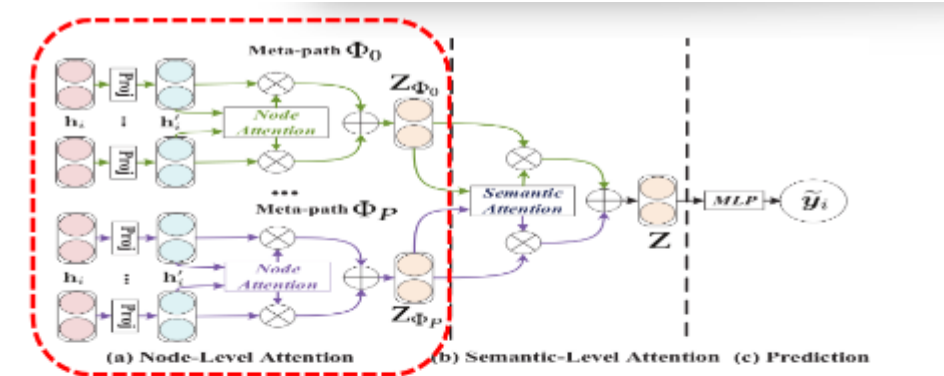


Model heterogeneous structure.

Capture rich semantics.

Task-specific loss.

## ➤ Node-level Attention and Aggregating



## ➤ Type-specific information

$$\mathbf{h}'_i = \mathbf{M}_{\phi_i} \cdot \mathbf{h}_i,$$

Type-specific transformation matrix

## ➤ Importance of Neighbors

$$e_{ij}^{\Phi} = att_{node}(\mathbf{h}'_i, \mathbf{h}'_j; \Phi).$$

$$\alpha_{ij}^{\Phi} = softmax_j(e_{ij}^{\Phi}) = \frac{\exp(\sigma(\mathbf{a}_{\Phi}^T \cdot [\mathbf{h}'_i \| \mathbf{h}'_j]))}{\sum_{k \in \mathcal{N}_i^{\Phi}} \exp(\sigma(\mathbf{a}_{\Phi}^T \cdot [\mathbf{h}'_i \| \mathbf{h}'_k]))},$$

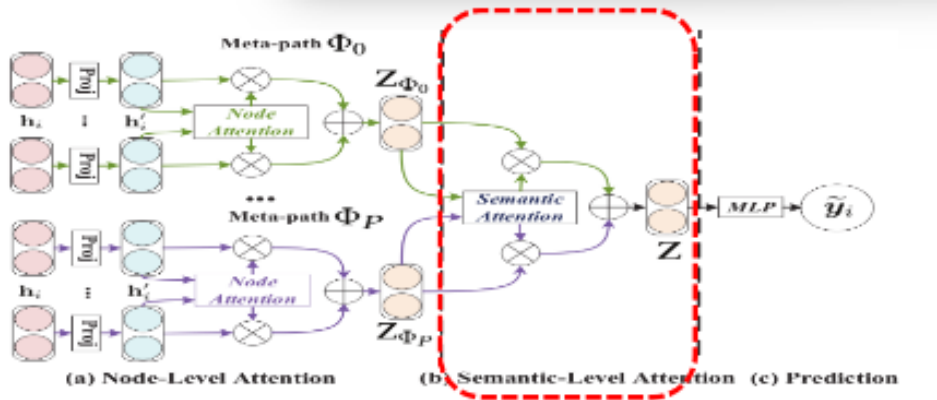
Node-level attention vector

## ➤ Node-level Aggregating

$$\mathbf{z}_i^{\Phi} = \bigoplus_{k=1}^K \sigma \left( \sum_{j \in \mathcal{N}_i^{\Phi}} \alpha_{ij}^{\Phi} \cdot \mathbf{h}'_j \right).$$

Node weight

## ➤ Semantic-level Attention and Aggregating



### ➤ Semantic-Level Attention

$$(\beta_{\Phi_0}, \beta_{\Phi_1}, \dots, \beta_{\Phi_P}) = att_{sem}(Z_{\Phi_0}, Z_{\Phi_1}, \dots, Z_{\Phi_P})$$

### ➤ Importance of Meta-path

Semantic-level attention vector


$$w_{\Phi_i} = \frac{1}{|\mathcal{V}|} \sum_{i \in \mathcal{V}} \mathbf{q}^T \cdot \tanh(\mathbf{W} \cdot \mathbf{z}_i^{\Phi} + \mathbf{b})$$

$$\beta_{\Phi_i} = \frac{\exp(w_{\Phi_i})}{\sum_{i=1}^P \exp(w_{\Phi_i})}$$

### ➤ Semantic-Level Aggregating

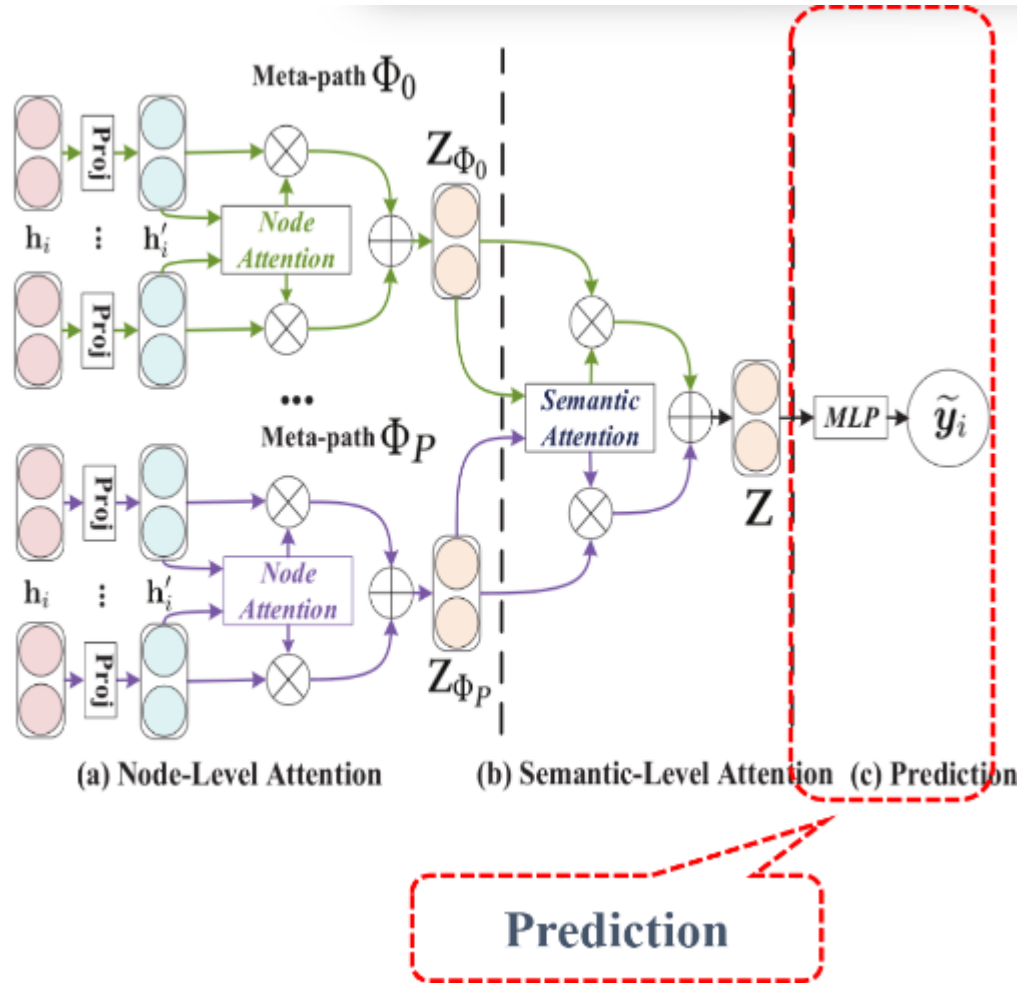
$$\mathbf{Z} = \sum_{i=1}^P \beta_{\Phi_i} \cdot \mathbf{Z}_{\Phi_i}$$

Semantic weight





## ➤ Prediction



## ➤ Semi-supervised Loss

$$L = - \sum_{l \in \mathcal{Y}_L} Y^l \ln(C \cdot Z^l)$$

Parameter of classifier

Labeled data

Optimize for the specific task  
(e.g. node classification)



- HAN is available as part of PyTorch Geometric library

```
from torch_geometric.nn import HANConv
```

```
class HAN(torch.nn.Module):
    def __init__(self, in_channels: Union[int, Dict[str, int]],
                  out_channels: int, hidden_channels=128, heads=8):
        super().__init__()
        self.han_conv = HANConv(in_channels, hidden_channels, heads=heads,
                                dropout=0.6, metadata=data.metadata())
        self.lin = Linear(hidden_channels, out_channels)

    def forward(self, x_dict, edge_index_dict):
        out = self.han_conv(x_dict, edge_index_dict)
        out = self.lin(out['movie'])
        return out
```



네트워크 과학연구실  
NETWORK SCIENCE LAB



가톨릭대학교  
THE CATHOLIC UNIVERSITY OF KOREA

